


Диффузионные модели




Что мы уже знаем


GAN: minimax the classification error loss.



VAE: maximize ELBO.




Flow-based generative models:
minimize the negative log-likelihood



Приложения




Stable Diffusion (StabilityAI)




Source: <https://learnopencv.com/image-generation-using-diffusion-models/>

Text-to-speech generation

InferGrad

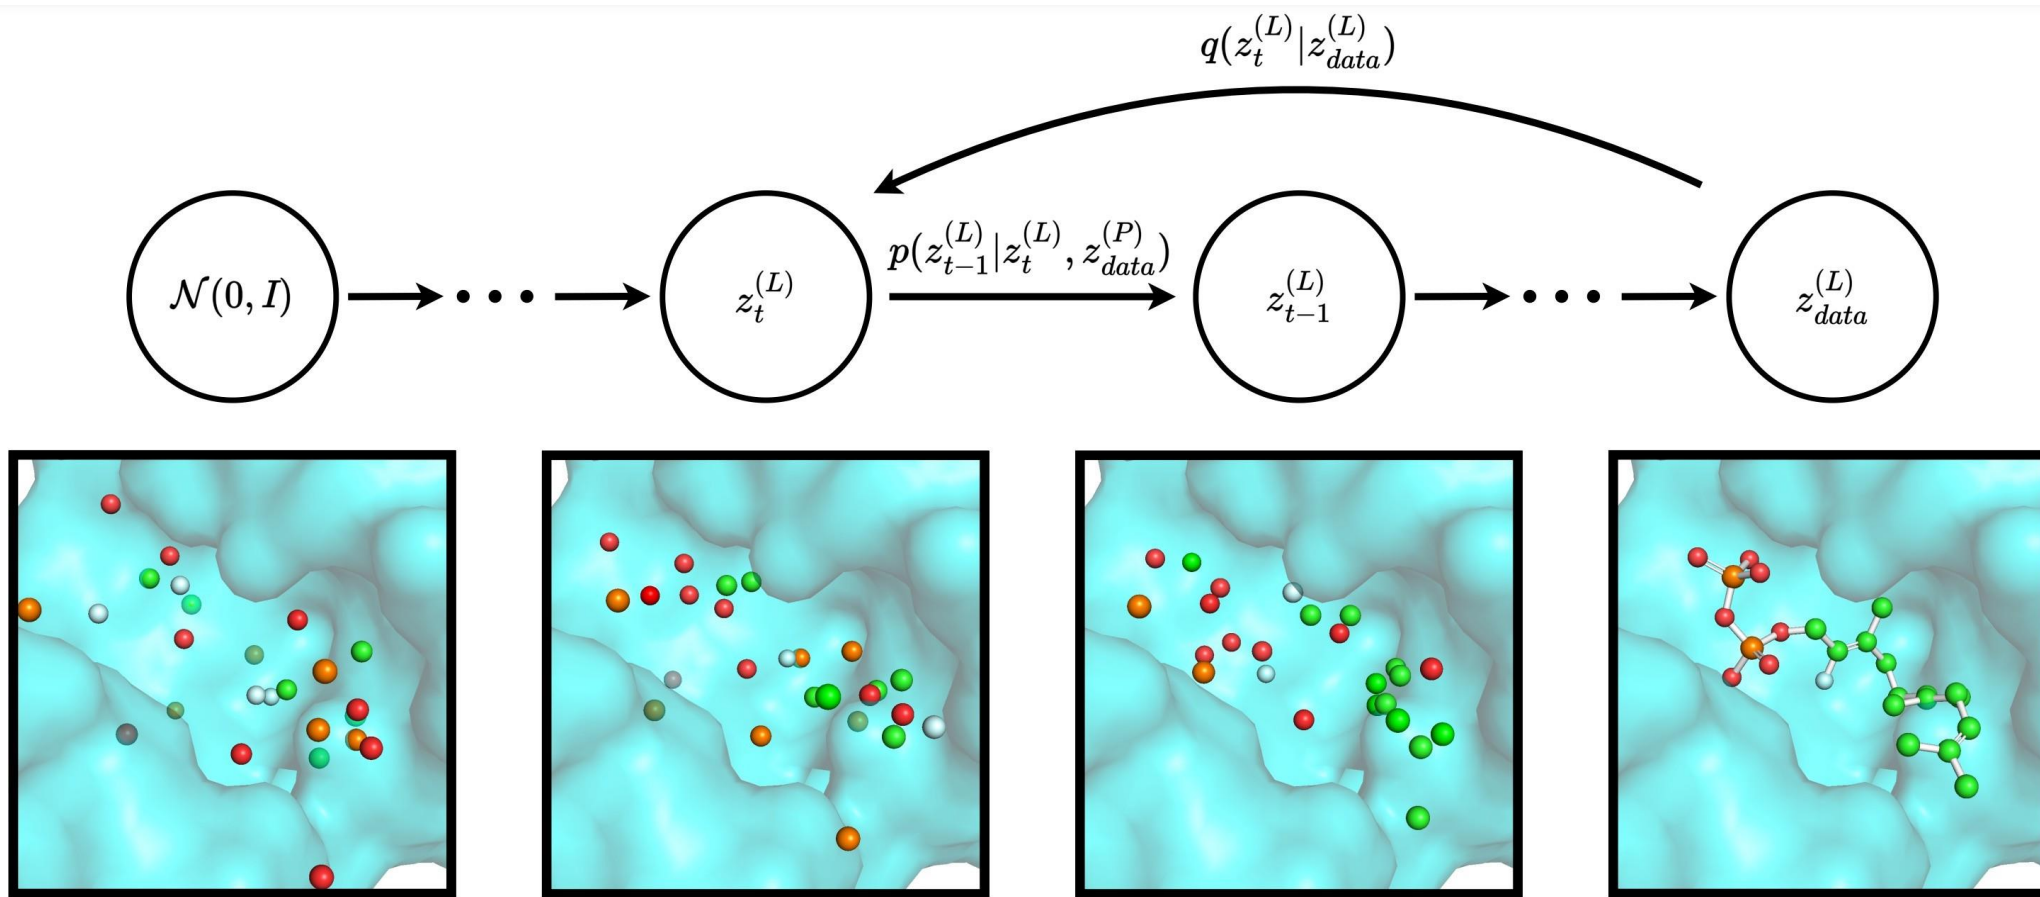


Ground Truth



Source: <https://github.com/heejkoo/Awesome-Diffusion-Models#text-to-speech>

Molecular and material generation



Source: <https://github.com/heejkoo/Awesome-Diffusion-Models#molecular-and-material-generation>

Astronomical spectra generation

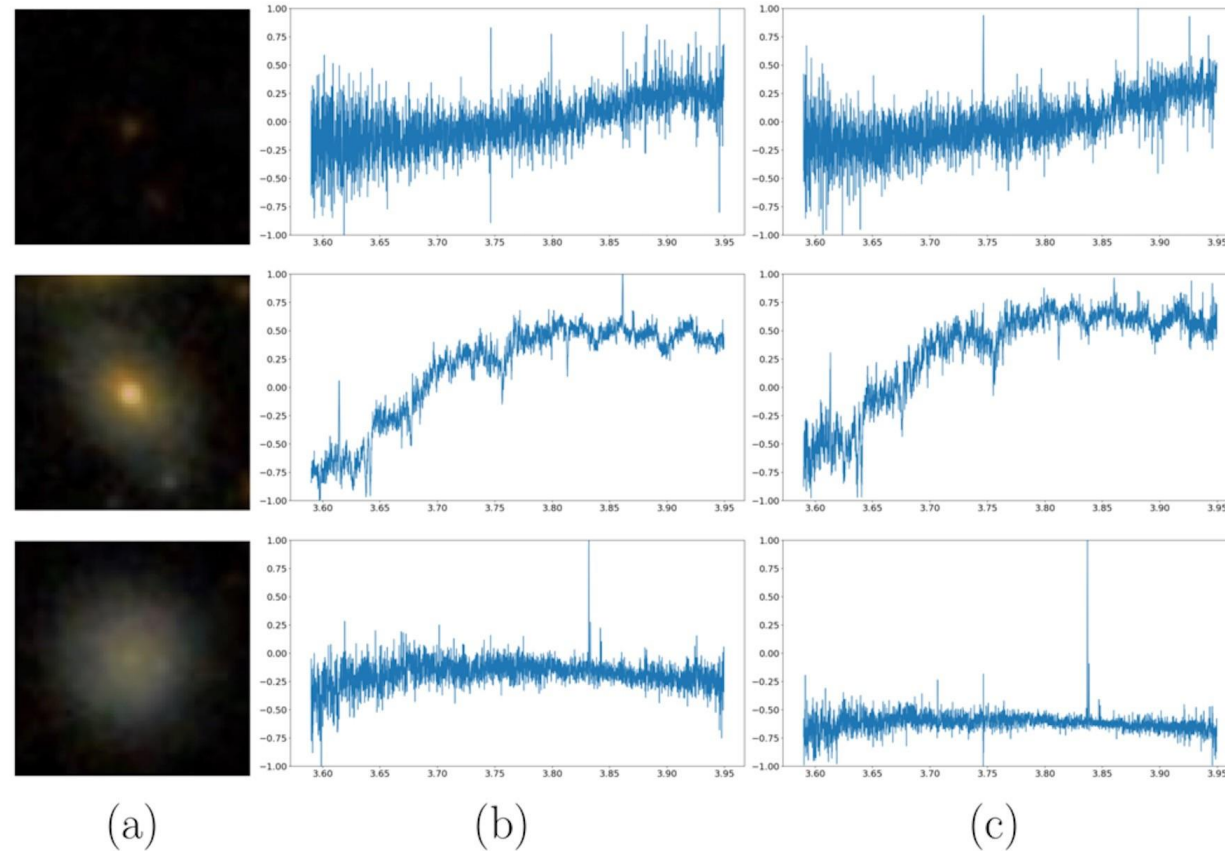


Figure 3: Generated spectra for the images in (a). In (b) we show the real spectra, in (c) the best match according to our contrastive model, out of 25 samples.

Source: https://ml4physicalsciences.github.io/2022/files/NeurIPS_ML4PS_2022_78.pdf

Dark matter density modelling




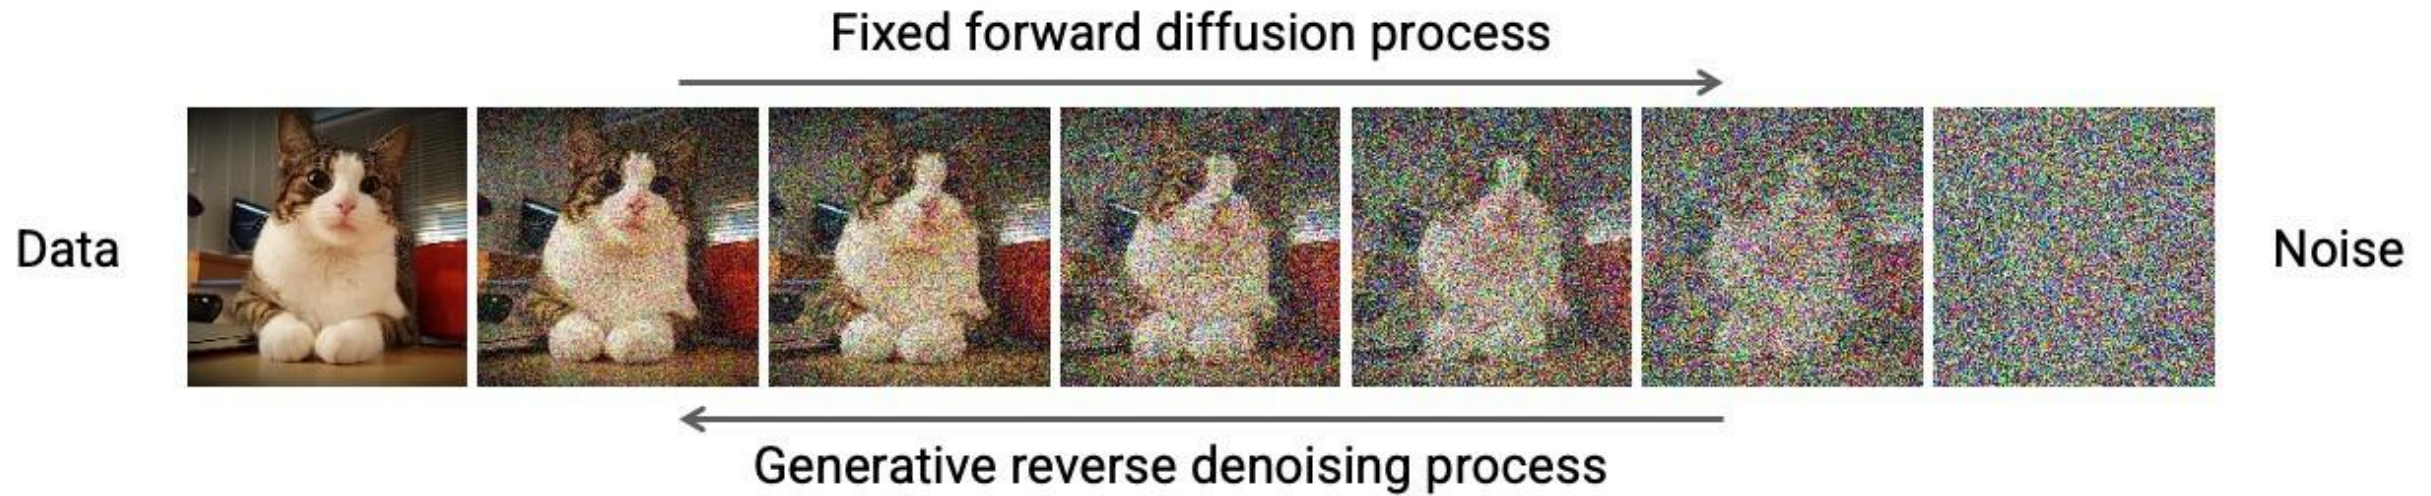
Figure 1: Four log cold dark matter mass density fields from the training data (top left) and from the sampled model (top right) at 128x128. Four samples of dust from the training data (bottom left) and from the trained model (bottom right).

Source: https://ml4physicalsciences.github.io/2022/files/NeurIPS_ML4PS_2022_25.pdf

Интуиция

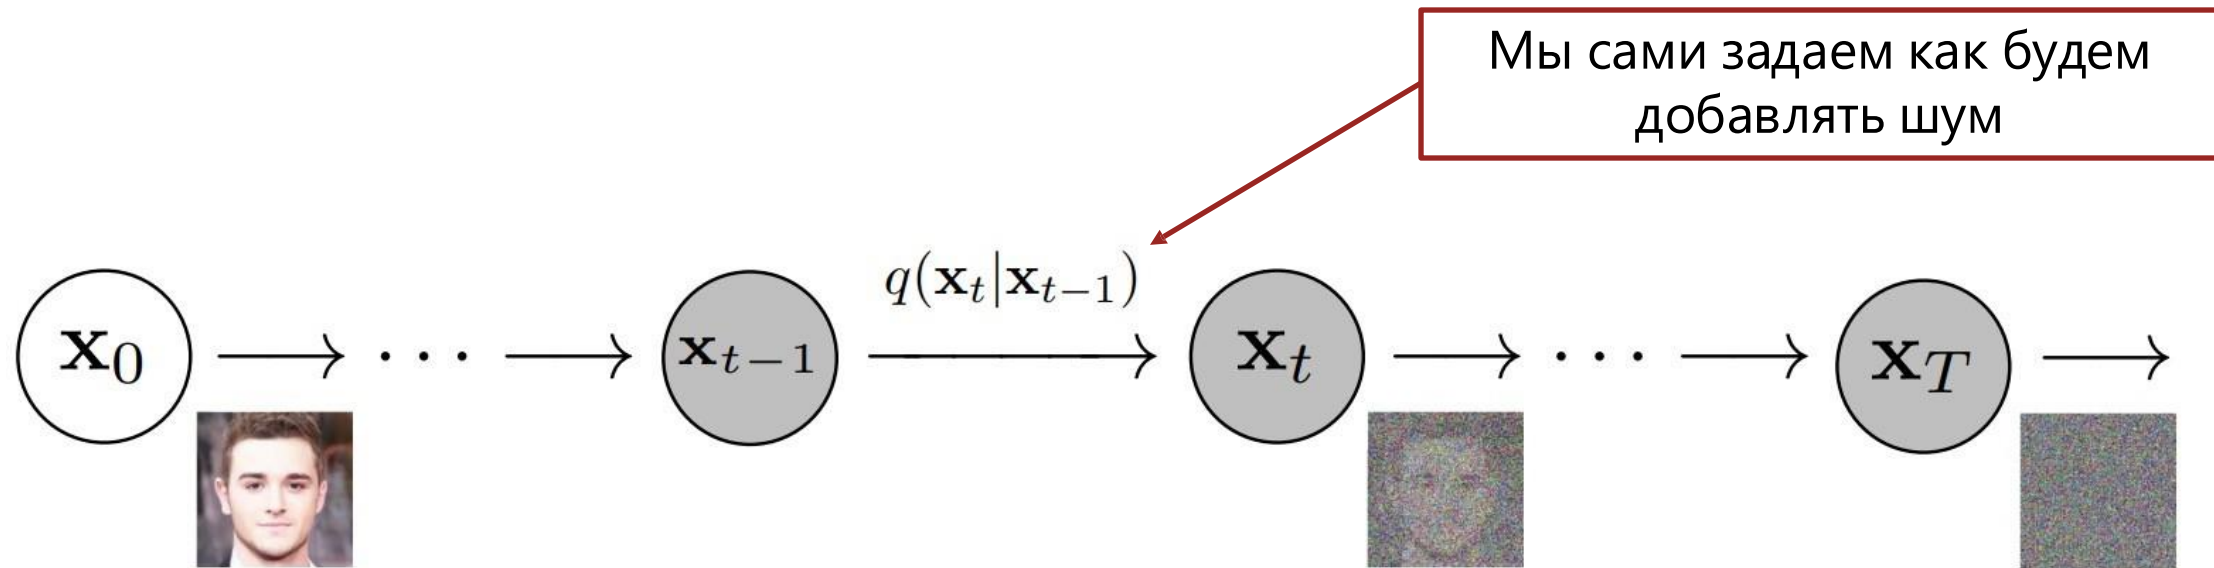


Общая идея диффузионных моделей



Источник: <https://cvpr2022-tutorial-diffusion-models.github.io/>

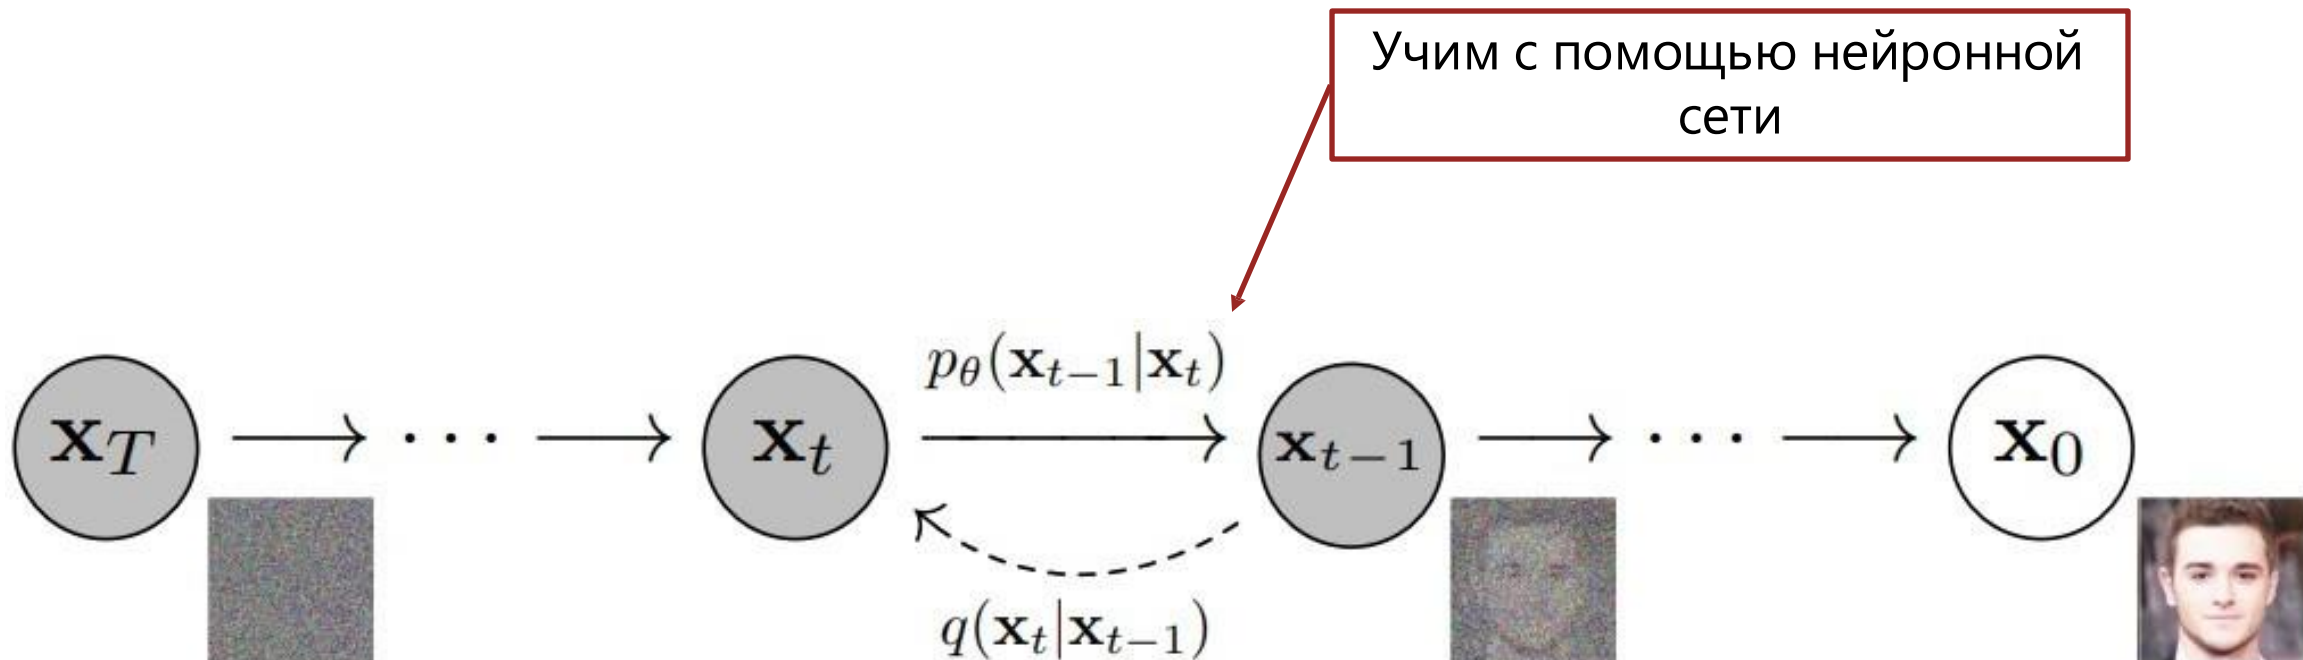
Процесс диффузии (зашумления)




$$q(\mathbf{x}_t | \mathbf{x}_{t-1}) = \mathcal{N}(\mathbf{x}_t; \boldsymbol{\mu}_t = \sqrt{1 - \beta_t} \mathbf{x}_{t-1}, \boldsymbol{\Sigma}_t = \beta_t \mathbf{I})$$

Константа

Обратный процесс




Демо: диффузия



Источник: <https://yang-song.net/blog/2021/score/>

Демо: обратный процесс




Источник: <https://yang-song.net/blog/2021/score/>

Диффузионные модели



Процесс диффузии (зашумления)




► $x_t = \sqrt{\alpha_t}x_{t-1} + \sqrt{1 - \alpha_t}\epsilon_t$

$$\epsilon_t \sim N(0, I), \quad \alpha_t = 1 - \beta_t$$

Мы сами так задали процесс

Процесс диффузии (зашумления)




- ▶ $x_t = \sqrt{\alpha_t}x_{t-1} + \sqrt{1 - \alpha_t}\epsilon_t$

$$\epsilon_t \sim N(0, I), \quad \alpha_t = 1 - \beta_t$$

- ▶ $x_t = \sqrt{\bar{\alpha}_t}x_0 + \sqrt{1 - \bar{\alpha}_t}\epsilon$

$$\epsilon \sim N(0, I), \quad \bar{\alpha}_t = \prod_{i=1}^t \alpha_i$$

Процесс диффузии (зашумления)



$$x_t = \sqrt{\alpha_t} x_{t-1} + \sqrt{1 - \alpha_t} \epsilon_t$$

$$\epsilon_t \sim N(0, I), \quad \alpha_t = 1 - \beta_t$$


$$x_t = \sqrt{\bar{\alpha}_t} x_0 + \sqrt{1 - \bar{\alpha}_t} \epsilon$$

$$\epsilon \sim N(0, I), \quad \bar{\alpha}_t = \prod_{i=1}^t \alpha_i$$

$$x_{t-1} = \frac{1}{\sqrt{\alpha_t}} \left(x_t - \frac{1 - \alpha_t}{\sqrt{1 - \bar{\alpha}_t}} \epsilon \right) + \tilde{\beta}_t z$$

$$z \sim N(0, I), \quad \tilde{\beta}_t = \frac{1 - \bar{\alpha}_{t-1}}{1 - \bar{\alpha}_t} \beta_t$$


Обратный процесс



► $\hat{x}_{t-1} = \frac{1}{\sqrt{\alpha_t}} \left(x_t - \frac{1-\alpha_t}{\sqrt{1-\bar{\alpha}_t}} \epsilon_\theta(x_t, t) \right) + \sigma_t z$ $z \sim N(0, I), \quad \sigma_t = const$

Предсказываем нейронной сетью

Обратный процесс



- ▶ $\hat{x}_{t-1} = \frac{1}{\sqrt{\alpha_t}} \left(x_t - \frac{1-\alpha_t}{\sqrt{1-\bar{\alpha}_t}} \epsilon_\theta(x_t, t) \right) + \sigma_t z$ $z \sim N(0, I), \quad \sigma_t = const$



- ▶ Функция потерь для обучения:

$$L_t = \|x_{t-1} - \hat{x}_{t-1}\|_2^2 \propto \|\epsilon - \epsilon_\theta(x_t, t)\|_2^2 \rightarrow \min_{\theta}$$

Архитектура нейронной сети



Time Representation



知乎 @华年ss

Источник: <https://www.zhihu.com/question/536012286/answer/2683123893>

Алгоритм обучения

Algorithm 1 Training

- 1: **repeat**
 - 2: $\mathbf{x}_0 \sim q(\mathbf{x}_0)$
 - 3: $t \sim \text{Uniform}(\{1, \dots, T\})$
 - 4: $\boldsymbol{\epsilon} \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$
 - 5: Take gradient descent step on
$$\nabla_{\theta} \left\| \boldsymbol{\epsilon} - \boldsymbol{\epsilon}_{\theta}(\sqrt{\bar{\alpha}_t} \mathbf{x}_0 + \sqrt{1 - \bar{\alpha}_t} \boldsymbol{\epsilon}, t) \right\|^2$$
 - 6: **until** converged
-

Алгоритм генерации

Algorithm 2 Sampling

```
1:  $\mathbf{x}_T \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$ 
2: for  $t = T, \dots, 1$  do
3:    $\mathbf{z} \sim \mathcal{N}(\mathbf{0}, \mathbf{I})$  if  $t > 1$ , else  $\mathbf{z} = \mathbf{0}$ 
4:    $\mathbf{x}_{t-1} = \frac{1}{\sqrt{\alpha_t}} \left( \mathbf{x}_t - \frac{1-\alpha_t}{\sqrt{1-\bar{\alpha}_t}} \boldsymbol{\epsilon}_{\theta}(\mathbf{x}_t, t) \right) + \sigma_t \mathbf{z}$ 
5: end for
6: return  $\mathbf{x}_0$ 
```


Значения гиперпараметров

- ▶ $T = 1000$
- ▶ $\beta_1 < \beta_2 < \dots < \beta_t < \dots < \beta_T$
- ▶ $\beta_1 = 0.0001, \beta_T = 0.02$
- ▶ $\sigma^2 = \beta_t$

Заключение




Generative learning trilemma




Source: <https://zhuanlan.zhihu.com/p/503932823>

Заключение


GAN: minimax the classification error loss.



VAE: maximize ELBO.



Flow-based generative models: minimize the negative log-likelihood



Diffusion models:
Gradually add Gaussian noise and then reverse

